

# AI Agent Architecture Document: PriceWarden

## System Overview

PriceWarden is an agentic AI system that automates e-commerce price comparison through intelligent reasoning, planning, and execution. The agent processes multimodal inputs (text/images) to find the best prices across Amazon, Flipkart, and Myntra.

## Core Components

### 1. PriceComparisonAgent (agent/core.py)

The central orchestrator implementing the agentic framework:

```
python

class PriceComparisonAgent:
    def __init__(self):
        self.vision_model = ImageToTextModel() # Fine-tuned BLIP
        self.scrapper = WebScraperTool()
        self.memory = SQLiteMemory()
        self.confidence_scorer = ConfidenceScorer()
```

### 2. Vision Model (models/inference.py)

- **Base Model:** Salesforce/blip-image-captioning-base
- **Fine-tuning:** LoRA adapter for e-commerce specialization
- **Purpose:** Converts product images to searchable text queries

### 3. Web Scraper Tool (scrapers/web\_scraper.py)

- **Dual-mode:** Requests for simple HTML, Selenium for JavaScript
- **Features:** Retry logic, rate limiting, fallback mechanisms
- **Sites:** Amazon, Flipkart, Myntra

### 4. Memory System (agent/memory.py)

- **Database:** SQLite for persistent storage
- **Tracks:** Search history, success rates, user preferences
- **Learning:** Improves confidence scores over time

## Interaction Flow

User Input (Text/Image)

↓

#### 1. REASONING PHASE

- Input type detection
- Image → Text conversion (if needed)
- Query refinement (remove stopwords)

↓

#### 2. PLANNING PHASE

- Site selection from config
- Priority ordering based on product type
- Resource allocation

↓

#### 3. EXECUTION PHASE

- Parallel web scraping
- Error handling & retries
- Fallback to cache if needed

↓

#### 4. LEARNING PHASE

- Store results in memory
- Update confidence scores
- Pattern recognition

↓

#### 5. SYNTHESIS PHASE

- Relevance scoring
- Price comparison
- Recommendation generation

↓

Final Output (Sorted Results + Insights)

## Detailed Component Interactions

### Input Processing Pipeline

python

```
def _reason(self, input_data):  
    if isinstance(input_data, Image):  
        # Use fine-tuned vision model  
        query = self.vision_model.generate_caption(input_data)  
        query = self._refine_search_query(query)  
    else:  
        # Direct text processing  
        query = self._refine_search_query(input_data)  
    return query
```

## Parallel Execution Strategy

python

```
def _execute(self, sites):
    with ThreadPoolExecutor(max_workers=3) as executor:
        futures = {
            executor.submit(self.scaper.scrape, site): site
            for site in sites
        }
        results = {}
        for future, site in futures.items():
            try:
                results[site] = future.result(timeout=10)
            except Exception as e:
                results[site] = self._get_fallback_data(site)
        return results
```

## Learning Mechanism

python

```
def _learn(self, query, results):
    # Update search history
    self.memory.add_search(query, results)

    # Calculate success metrics
    success_rate = len(results) / len(self.sites)
    avg_relevance = np.mean([r.relevance for r in results])

    # Update confidence
    self.confidence_scorer.update(success_rate, avg_relevance)

    # Pattern recognition
    if "kurta" in query.lower():
        self.memory.update_site_preference("mynta", +0.1)
```

## Technology Choices & Rationale

### Why BLIP + LoRA?

- **BLIP**: Strong baseline for image captioning
- **LoRA**: Enables fine-tuning on consumer GPU (8GB VRAM)
- **Result**: 73% accuracy with 99.75% fewer trainable parameters

### Why Selenium + Requests?

- **Requests:** Fast for static content (Amazon product lists)
- **Selenium:** Necessary for JavaScript-rendered prices
- **Fallback:** Ensures reliability even when scraping fails

## Why SQLite?

- **Lightweight:** No server setup required
- **Sufficient:** Handles our data volume easily
- **Portable:** Easy to deploy and backup

## Why ThreadPoolExecutor?

- **Simplicity:** Easier than asyncio for this use case
- **Performance:** 3x speedup (15s → 5s)
- **Control:** Better error handling than multiprocessing

## Configuration Management

All settings centralized in `config/config.yaml`:

```
yaml
```

```
agent:
  confidence_threshold: 0.7
  max_retries: 3
  timeout: 10

model:
  checkpoint_path: "./models/checkpoints/final_model"
  device: "cuda"
  max_length: 50

scraping:
  sites:
    - name: "amazon"
      url: "https://www.amazon.in/s?k="
      use_selenium: false
    - name: "flipkart"
      url: "https://www.flipkart.com/search?q="
      use_selenium: true
    - name: "myntra"
      url: "https://www.myntra.com/"
      use_selenium: true
      fallback_enabled: true

rate_limiting:
  requests_per_minute: 30
  delay_between_sites: 1
```

## Error Handling Strategy

### Multi-Layer Resilience

1. Try primary scraping method
2. On failure → Retry with exponential backoff
3. Still failing → Switch to Selenium
4. Persistent failure → Use cached fallback data
5. Log all failures for debugging

### Example Implementation

```
python
```

```
def scrape_with_resilience(self, url):
    strategies = [
        self._scrape_with_requests,
        self._scrape_with_selenium,
        self._get_cached_data
    ]

    for strategy in strategies:
        try:
            return strategy(url)
        except Exception as e:
            logger.warning(f"Strategy {strategy.__name__} failed: {e}")
            continue

    return [] # Empty results better than crashing
```

## Performance Optimizations

### Model Loading

- Lazy loading: Model loads only when needed
- Cached in memory after first load
- Quantized to INT8 for 2x faster inference

### Scraping Optimization

- Connection pooling for HTTP requests
- Pre-compiled regex patterns
- CSS selector caching

### Memory Management

- Limited to 100 products in memory
- Automatic cleanup of old searches (>24 hours)
- Efficient data structures (only store essentials)

## Scalability Considerations

### Horizontal Scaling Ready

- Stateless agent design
- Database supports concurrent access
- Could deploy multiple instances behind load balancer

## Vertical Scaling Options

- Batch processing for multiple queries
- GPU inference for higher throughput
- Caching layer (Redis) for common searches

## Security & Ethics

### Web Scraping Ethics

- Respects robots.txt
- Rate limiting (30 requests/minute)
- User-agent identification
- No personal data collection

### User Privacy

- All processing done locally
- No user data stored permanently
- Anonymous usage statistics only

## System Requirements

### Minimum Requirements

- Python 3.8+
- 8GB RAM
- 4GB GPU VRAM (or CPU mode)
- 2GB disk space

### Recommended Setup

- Python 3.10
- 16GB RAM
- 8GB GPU VRAM
- SSD for database

## Deployment Architecture

User Interface (Streamlit)



CostFilter Agent



Vision Scraper Memory		
Model	Tool	System
GPU	Internet	SQLite

## Monitoring & Logging

### Metrics Tracked

- Query response time
- Scraping success rate
- Model inference speed
- Memory usage
- Cache hit rate

### Logging Levels

- INFO: Normal operations
- WARNING: Fallback activations
- ERROR: Scraping failures
- DEBUG: Detailed execution flow

## Future Architecture Enhancements

### Multi-Agent System

Orchestrator Agent	
├──	Search Agent (current functionality)
├──	Price Tracking Agent (monitors changes)
├──	Recommendation Agent (personalized suggestions)
└──	Negotiation Agent (finds coupons/deals)

### Microservices Architecture

- Vision Service (Docker container)
- Scraping Service (distributed workers)
- API Gateway (FastAPI)
- Database Service (PostgreSQL upgrade)